


MiniMind medical LoRA fine-tuning

64M base model | 5,000 training examples | 1 epoch | rank 16

Batch size 8 | Sequence length 512 | 100 held-out validation examples



Single run (seed 42); 14 measured checkpoints connected by straight lines; no smoothing.

The y-axis is zoomed. Loss reduction does not establish an improvement in medical accuracy.