


Validation loss keeps improving after accuracy stops



Lower loss on the supervised tokens no longer converts into more correct classifications — the ceiling is model capacity, not data.